

MCP's default costs 10-30x more than it needs to.

Two independent sources landed on
the same number this year.

Anthropic's own engineering team found this.

150,000 → 2,000 tokens.

98.7% cut, one real example.

The fix: have the agent write code against MCP servers exposed as a filesystem, instead of loading every tool schema and result through itself.

A third party benchmarked it independently. Same result.

CLI, PER TASK

1.3k-9.4k

MCP, PER TASK

32k-82k tokens

~\$3.20/mo → ~\$55.20/mo

projected cost, 10k ops/month.

Reliability took a hit too.



MCP: 72% success — TCP timeouts.



CLI: 100% success.

Same 5 real GitHub tasks,
same agent, same model.

Every call injects the full tool schema list — whether the task needs it or not.

Every intermediate result routes back through the model instead of being handled directly.

I've spent time this year evaluating MCP-based tooling at work.

The protocol isn't the problem. Eager-loading everything by default is.

If you're running MCP in prod - loading full schemas every call, or lazy-loading? What did the eager version cost you?