

Postgres isn't slow. Your storage is.

Most 'Postgres is slow' complaints
trace back to one layer nobody blames.


A POSETTE 2026 talk made the case directly.

P95 spikes. Autovacuum falling behind.
Checkpoint stalls. Replication lag.
The usual complaint list -
mostly one layer underneath it all.

Sai Srirampur, POSETTE 2026

recapped by ClickHouse, 1 Sep 2026

THE SAME WRITE, TWO STORAGE LAYERS



Same workload. Same query. Different floor under it.

WHAT ACTUALLY HAPPENS



That reframes the debugging question:
not 'why is Postgres slow,' but
'what's under my WAL writes.'

I haven't run this comparison myself.

No local-NVMe box to test against a managed network-storage instance side by side.

But fsync and checkpoint I/O are exactly the operations most sensitive to storage latency — the mechanism holds without needing to re-run it.

Anyone debugging Postgres on managed cloud storage.



**Cloud-managed
Postgres, default**



**Checkpoint tuning
before storage swap**



**Benchmarking the
wrong layer**

Most teams tune Postgres knobs
before ever checking what's
underneath the WAL.

If you've moved a Postgres workload between storage layers -

**What actually changed,
and was it what you expected?**

**Network block storage vs. local NVMe -
same workload, different floor.**