

**p95 latency crept
up on a platform
moving 10M+ msgs
a month.**

Here's what I actually found —
and what I still don't have a
good answer for.

p95 latency:

The response time 95% of requests are faster than.

Not the average — averages hide the worst realistic case. p95 is closer to what your slowest real users actually feel.

I assumed it was network overhead between services.

It wasn't.

A hot read endpoint was hitting the database directly — for data that barely changed request to request.

THE FIX

Redis cache-aside on that one query.

RESULT

↓ ~20%

p95 latency, on that endpoint.

**What I still don't
have a clean
answer for:**

**How to size the cache TTL so
staleness never shows up to a
user — without just picking a
number and hoping.**